

Computational Statistics - Lab 04

Annalena Erhard (anner218) and Maximilian Pfundstein (maxpf364)

2019-02-04

Contents

1	Question 1: Computations with Metropolis-Hastings	1
2	Question 2: Gibbs Sampling	4
3	Source Code	5

1 Question 1: Computations with Metropolis-Hastings

Consider the following probability density function:

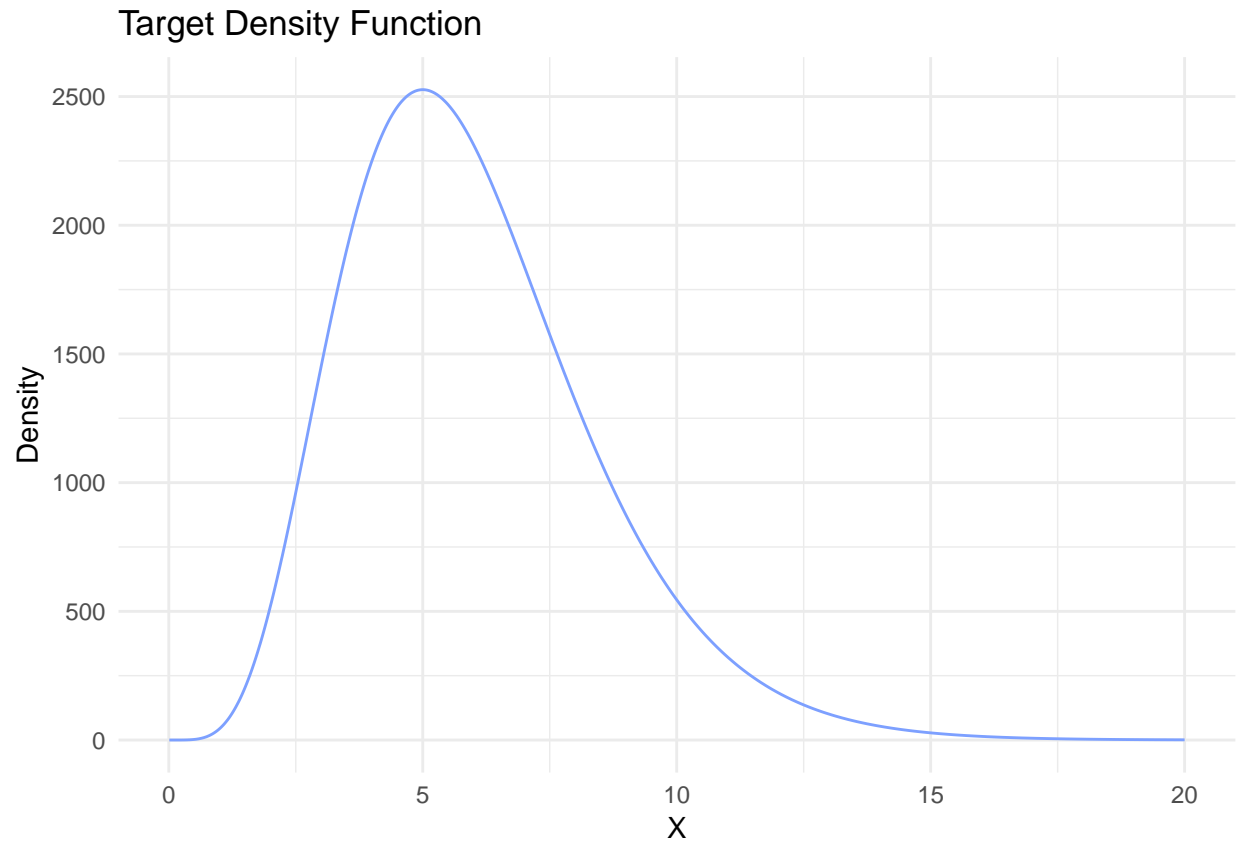
$$f(x) \propto x^5 e^{-x}, \quad x > 0.$$

You can see that the distribution is known up to some constant of proportionality. If you are interested (**NOT** part of the Lab) this constant can be found by applying integration by parts multiple times and equals 120.

1. Use Metropolis-Hastings algorithm to generate samples from this distribution by using proposal distribution as log-normal $LN(X_t, 1)$, take some starting point. Plot the chain you obtained as a time series plot. What can you guess about the convergence of the chain? If there is a burn-in period, what can be the size of this period?

```
# Target function with original scaling
f = function(x) {
  return(120 * x^5 * exp(-x))
}

# Target function
df = function(x) {
  if (sum(c(x) <= 0) > 0) stop("x must be larger than 0.")
  return(x^5 * exp(-x))
}
```



```
## Metropolis Hasting Algorithm
##
## @param n Number of samples from the target distribution.
## @param x_0 Initial state from Pi.
## @param b Burn-in steps before taking samples.
##
## @return Returns a list containing the burned and normal samples.
## @export
##
## @examples
metropolis_hasting = function(n = 1, x_0 = 1, b = 50) {

  # Vectors to store the samples in.
  burn_in_samples = c()
  samples = c()

  # Samples from proposal from the random walk (MC)
  xt = x_0
  xt_1 = x_0

  walk = function(burn) {
    # Generate proposal state
    x_star = rlnorm(n = 1, meanlog = xt, sdlog = 1)

    # Calculate correction factor C
    c = dlnorm(xt_1, meanlog = xt, sdlog = 1) /
```

```

    dlnorm(x_star, meanlog = xt, sdlog = 1)

    # Calculate acceptance probability alpha
    alpha = min(1, df(x_star)/df(xt_1))

    # Generate u from uniform
    u = runif(n = 1, min = 0, max = 1)

    # Decide if to accept or reject the proposal
    if (u <= alpha) {
      # Accept
      xt_1 <- xt
      xt <- x_star

      # Depending on if we're still burning save in different vector
      if (burn) {
        burn_in_samples <- c(burn_in_samples, x_star)
      }
      else {
        samples <- c(samples, x_star)
      }
    }
    else {
      # Reject
      xt <- xt_1
    }
  }


  # Burn-in period
  while (length(burn_in_samples) < n) {
    walk(TRUE)
  }

  # Draw real samples
  while (length(samples) < n) {
    walk(FALSE)
  }

  # Return the burned and normal samples
  return(list(samples = samples, burn_in_samples = burn_in_samples))
}

results = metropolis_hasting(10000, x_0 = 5, b = 100)

```



2. Perform Step 1 by using the chi-square distribution $\chi^2(\lfloor X_t + 1 \rfloor)$ as a proposal distribution, where $\lfloor x \rfloor$ is the floor function, meaning the integer part of x for positive x , i.e. $\lfloor 2.95 \rfloor = 2$.
3. Compare the results of Steps 1 and 2 and make conclusions.
4. Generate 10 MCMC sequences using the generator from Step 2 and starting points 1, 2, ..., or 10. Use the Gelman–Rubin method to analyze convergence of these sequences.
5. Estimate

$$\int_0^{\infty} x f(x) dx$$

using the samples from Steps 1 and 2.

6. The distribution generated is in fact a gamma distribution. Look in the literature and define the actual value of the integral. Compare it with the one you obtained.

2 Question 2: Gibbs Sampling

concentration of a certain chemical was measured in a water sample, and the result was stored in the data `chemical.RData` having the following variables:

- X : day of the measurement
- Y : measured concentration of the chemical.

The instrument used to measure the concentration had certain accuracy; this is why the measurements can be treated as noisy. Your purpose is to restore the expected concentration values.

1. Import the data to R and plot the dependence of Y on X. What kind of model is reasonable to use here?
2. A researcher has decided to use the following (random-walk) Bayesian model (n=number of observations, $\vec{\mu} = (\mu_1, \dots, \mu_n)$ are unknown parameters):

$$Y_i = \mathcal{N}(\mu_i, \text{variance} = 0.2), \quad i = 1, \dots, n$$

where the prior is

$$p(\mu_1) = 1$$

$$p(\mu_{i+1}|\mu_i) = \mathcal{N}(\mu_i, 0.2), i = 1, \dots, n-1.$$

Present the formulae showing the likelihood $p(\vec{Y}|\vec{\mu})$ and the prior $p(\vec{\mu})$. **Hint:** a chain rule can be used here $p(\vec{\mu}) = p(\mu_1)p(\mu_2|\mu_1)p(\mu_3|\mu_2)\dots p(\mu_n|\mu_{n-1})$

3. Use Bayes' Theorem to get the posterior up to a constant proportionality, and then find out the distributions of $(\mu_i|\vec{\mu}_{-i}, \vec{Y})$, where μ_{-i} is a vector containing all μ values except of μ_i .

Hint A: consider for separate formulae for $(\mu_1|\vec{\mu}_{-1}, \vec{Y})$, $(\mu_n|\vec{\mu}_{-n}, \vec{Y})$ and then a formula for all remaining $(\mu_i|\vec{\mu}_{-i}, \vec{Y})$.

Hint B:

$$e^{-\frac{1}{d}((x-a)^2+(x-b)^2)} \propto e^{-\frac{(x-(a+b)/2)^2}{d/2}}$$

Hint C:

$$e^{-\frac{1}{d}((x-a)^2+(x-b)^2+(x-c)^2)} \propto e^{-\frac{(x-(a+b+c)/3)^2}{d/3}}$$

4. Use the distributions derived in Step 3 to implement a Gibbs sampler that uses $\vec{\mu}^0 = 0, \dots, 0$ as a starting point. Run the Gibbs sampler to obtain 1000 values of $\vec{\mu}$ and then compute the expected value of $\vec{\mu}$ versus X and Y versus X in the same graph. Does it seem that you have managed to remove the noise? Does it seem that the expected value of $\vec{\mu}$ can catch the true underlying dependence between Y and X?
5. Make a trace plot for μ_n and comment on the burn-in period and convergence.

```
# Loading RData
#data2 = get(load("data.RData"))
#head(data2)
```

3 Source Code

```
knitr::opts_chunk$set(echo = TRUE, cache = FALSE, include = TRUE, eval = TRUE)
library(knitr)
library(readxl)
library(ggplot2)
library(gridExtra)

# Target function with original scaling
f = function(x) {
```

```

    return(120 * x^5 * exp(-x))
}

# Target function
df = function(x) {
  if (sum(c(x) <= 0) > 0) stop("x must be larger than 0.")
  return(x^5 * exp(-x))
}

sequence = seq(from = 0.01, to = 20, by = 0.01)

real_f = f(sequence)
plotdf = data.frame(sequence, real_f)

ggplot(plotdf) +
  geom_line(aes(x = sequence, y = real_f), color = "#7da0ff") +
  labs(title = "Target Density Function", y = "Density",
       x = "X", color = "Legend") +
  theme_minimal()

#' Metropolis Hasting Algorithm
#'
#' @param n Number of samples from the target distribution.
#' @param x_0 Initial state from Pi.
#' @param b Burn-in steps before taking samples.
#'
#' @return Returns a list containing the burned and normal samples.
#' @export
#'
#' @examples
metropolis_hasting = function(n = 1, x_0 = 1, b = 50) {

  # Vectors to store the samples in.
  burn_in_samples = c()
  samples = c()

  # Samples from proposal from the random walk (MC)
  xt = x_0
  xt_1 = x_0

  walk = function(burn) {
    # Generate proposal state
    x_star = rlnorm(n = 1, meanlog = xt, sdlog = 1)

    # Calculate correction factor C
    c = dlnorm(xt_1, meanlog = xt, sdlog = 1) /
      dlnorm(x_star, meanlog = xt, sdlog = 1)

    # Calculate acceptance probability alpha
    alpha = min(1, df(x_star)/df(xt_1))
  }
}

```

```

# Generate u from uniform
u = runif(n = 1, min = 0, max = 1)

# Decide if to accept or reject the proposal
if (u <= alpha) {
  # Accept
  xt_1 <- xt
  xt <- x_star

  # Depending on if we're still burning save in different vector
  if (burn) {
    burn_in_samples <- c(burn_in_samples, x_star)
  }
  else {
    samples <- c(samples, x_star)
  }
}
else {
  # Reject
  xt <- xt_1
}
}

# Burn-in period
while (length(burn_in_samples) < n) {
  walk(TRUE)
}

# Draw real samples
while (length(samples) < n) {
  walk(FALSE)
}

# Return the burned and normal samples
return(list(samples = samples, burn_in_samples = burn_in_samples))
}

results = metropolis_hasting(10000, x_0 = 5, b = 100)

plotdf = data.frame(results$samples)

ggplot(plotdf) +
  geom_histogram(aes(x = results.samples,
                    color = "#000000", fill = "#C70039", bins = length(results$samples)/100) +
  labs(title = "Samples From Target Function", y = "Density",
        x = "X", color = "Legend") +
  theme_minimal()

# Loading RData
#data2 = get(load("data.RData"))
#head(data2)

```